

shatery

January 16, 2024

#Bank Loan Model Project This project, I aim to develop a classifier for identifying potential customers who are more likely to purchase a personal loan using the Thera-Bank dataset. Thera-Bank is interested in expanding their loan business by converting liability customers into retail loan customers, while keeping them as depositors. The retail marketing department is developing campaigns with better target marketing to increase the success rate with a minimal budget.

##content Column descriptions ID Customer ID Age Customer's age in completed years Experience #years of professional experience Income Annual income of the customer (\$000).

ZIP Code Home Address. Family Family size of the customer CCAvg Avg. spending on credit cards per month (\$) Education Level.

1: Undergrad; 2: Graduate; 3: Advanced/Professional Mortgage Value of house mortgage if any. (\$000) Personal Loan Did this customer accept the personal loan offered in the last campaign? Securities Account Does the customer have a securities account with the bank? CD Account Does the customer have a certificate of deposit (CD) account with the bank? Online Does the customer use internet banking facilities? CreditCard Does the customer uses a credit card issued by UniversalBank?

ID (Customer ID): Customer identification code.

Age: Age of the customer in completed years.

Experience: Number of years of professional experience.

Income: Annual income of the customer in thousands of dollars.

ZIPCode: Home address ZIP code.

Family: Family size of the customer.

CCAvg (Credit Card Average): Average spending on credit cards per month in thousands of dollars.

Education: Education level of the customer. (1: Undergraduate; 2: Graduate; 3: Advanced/Professional)

Mortgage: Value of the house mortgage, if any, in thousands of dollars.

Personal Loan: Whether the customer accepted the personal loan offered in the last campaign (Binary: 1 or 0).

Securities Account: Whether the customer has a securities account with the bank (Binary: 1 or 0).

CD Account: Whether the customer has a certificate of deposit (CD) account with the bank (Binary: 1 or 0).

Online: Whether the customer uses internet banking facilities (Binary: 1 or 0).

CreditCard: Whether the customer uses a credit card issued by UniversalBank (Binary: 1 or 0).

```
[1]: #Import library
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib.colors import ListedColormap
from sklearn import model_selection
from sklearn.naive_bayes import GaussianNB
from sklearn import metrics
from sklearn.feature_selection import mutual_info_classif
```

```
[2]: Bank_Personal = pd.read_excel('Bank_Personal_Loan_Modelling.xlsx',
    ↪sheet_name='Data')
```

#Get Some Information

```
[3]: Bank_Personal
```

```
[3]:
```

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	Education	\
0	1	25	1	49	91107	4	1.6	1	
1	2	45	19	34	90089	3	1.5	1	
2	3	39	15	11	94720	1	1.0	1	
3	4	35	9	100	94112	1	2.7	2	
4	5	35	8	45	91330	4	1.0	2	
...	
4995	4996	29	3	40	92697	1	1.9	3	
4996	4997	30	4	15	92037	4	0.4	1	
4997	4998	63	39	24	93023	2	0.3	3	
4998	4999	65	40	49	90034	3	0.5	2	
4999	5000	28	4	83	92612	3	0.8	1	

	Mortgage	Personal Loan	Securities Account	CD Account	Online	\
0	0	0	1	0	0	
1	0	0	1	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	
...	
4995	0	0	0	0	1	
4996	85	0	0	0	1	
4997	0	0	0	0	0	
4998	0	0	0	0	1	
4999	0	0	0	0	1	

	CreditCard
0	0
1	0
2	0
3	0
4	1
...	...
4995	0
4996	0
4997	0
4998	0
4999	1

[5000 rows x 14 columns]

```
[4]: Bank_Personal.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    5000 non-null   int64
1   Age                  5000 non-null   int64
2   Experience            5000 non-null   int64
3   Income               5000 non-null   int64
4   ZIP Code             5000 non-null   int64
5   Family               5000 non-null   int64
6   CCAvg               5000 non-null   float64
7   Education            5000 non-null   int64
8   Mortgage            5000 non-null   int64
9   Personal Loan        5000 non-null   int64
10  Securities Account    5000 non-null   int64
11  CD Account           5000 non-null   int64
12  Online               5000 non-null   int64
13  CreditCard           5000 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 547.0 KB
```

```
[5]: Bank_Personal.apply(lambda x: len(x.unique()))
```

```
[5]: ID                    5000
     Age                    45
     Experience            47
     Income                162
     ZIP Code             467
     Family                 4
```

```

CCAvg          108
Education      3
Mortgage       347
Personal Loan   2
Securities Account  2
CD Account     2
Online         2
CreditCard     2
dtype: int64

```

```
[6]: Bank_Personal.describe()
```

```

[6]:
      count      ID      Age  Experience      Income      ZIP Code  \
count  5000.000000  5000.000000  5000.000000  5000.000000  5000.000000
mean    2500.500000   45.338400   20.104600   73.774200  93152.503000
std     1443.520003   11.463166   11.467954   46.033729   2121.852197
min         1.000000   23.000000   -3.000000    8.000000   9307.000000
25%     1250.750000   35.000000   10.000000   39.000000   91911.000000
50%     2500.500000   45.000000   20.000000   64.000000   93437.000000
75%     3750.250000   55.000000   30.000000   98.000000   94608.000000
max      5000.000000   67.000000   43.000000  224.000000  96651.000000

      count      Family      CCAvg      Education      Mortgage      Personal Loan  \
count  5000.000000  5000.000000  5000.000000  5000.000000  5000.000000
mean     2.396400    1.937913    1.881000    56.498800    0.096000
std     1.147663    1.747666    0.839869   101.713802    0.294621
min     1.000000    0.000000    1.000000    0.000000    0.000000
25%     1.000000    0.700000    1.000000    0.000000    0.000000
50%     2.000000    1.500000    2.000000    0.000000    0.000000
75%     3.000000    2.500000    3.000000   101.000000    0.000000
max     4.000000   10.000000    3.000000   635.000000    1.000000

      count      Securities Account      CD Account      Online      CreditCard
count    5000.000000  5000.000000  5000.000000  5000.000000  5000.000000
mean         0.104400    0.06040    0.596800    0.294000
std         0.305809    0.23825    0.490589    0.455637
min         0.000000    0.00000    0.000000    0.000000
25%         0.000000    0.00000    0.000000    0.000000
50%         0.000000    0.00000    1.000000    0.000000
75%         0.000000    0.00000    1.000000    1.000000
max         1.000000    1.00000    1.000000    1.000000

```

```

[7]: #Number of customers
personal_loan_counts = Bank_Personal["Personal Loan"].value_counts()
print(f"Number of customers with Personal Loan:\n{personal_loan_counts[1]}")
print(f"Number of customers without Personal Loan:\n{personal_loan_counts[0]}")

```

Number of customers with Personal Loan:
480
Number of customers without Personal Loan:
4520

```
[8]: Bank_Personal.isnull().sum()
```

```
[8]: ID          0  
     Age         0  
     Experience  0  
     Income      0  
     ZIP Code    0  
     Family      0  
     CCAvg       0  
     Education   0  
     Mortgage    0  
     Personal Loan 0  
     Securities Account 0  
     CD Account  0  
     Online      0  
     CreditCard  0  
     dtype: int64
```

```
[9]: categorical_variables=[col for col in Bank_Personal.columns if  
    ↪Bank_Personal[col].nunique()<=5]  
print(categorical_variables)  
continuous_variables=[col for col in Bank_Personal.columns if  
    ↪Bank_Personal[col].nunique()>5]  
print(continuous_variables)
```

```
['Family', 'Education', 'Personal Loan', 'Securities Account', 'CD Account',  
'Online', 'CreditCard']  
['ID', 'Age', 'Experience', 'Income', 'ZIP Code', 'CCAvg', 'Mortgage']
```

```
[10]: #check duplicated rows  
Bank_Personal.duplicated().sum()
```

```
[10]: 0
```

#Data Preprocessing

```
[11]: # Change the data type of 'ZIP Code' to object  
Bank_Personal['ZIP Code'] = Bank_Personal['ZIP Code'].astype(object)  
  
#Change the data type of multiple columns to boolean  
columns_to_convert = ['CreditCard', 'Personal Loan', 'Securities Account', 'CD_  
    ↪Account', 'Online']
```

```
Bank_Personal[columns_to_convert] = Bank_Personal[columns_to_convert].
↳astype(bool)
```

```
[12]: Bank_Personal.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    5000 non-null   int64
1   Age                   5000 non-null   int64
2   Experience             5000 non-null   int64
3   Income                5000 non-null   int64
4   ZIP Code              5000 non-null   object
5   Family                5000 non-null   int64
6   CCAvg                 5000 non-null   float64
7   Education              5000 non-null   int64
8   Mortgage              5000 non-null   int64
9   Personal Loan         5000 non-null   bool
10  Securities Account    5000 non-null   bool
11  CD Account            5000 non-null   bool
12  Online                5000 non-null   bool
13  CreditCard            5000 non-null   bool
dtypes: bool(5), float64(1), int64(7), object(1)
memory usage: 376.1+ KB
```

```
[13]: # show number of Experience < 0
Bank_Personal[Bank_Personal['Experience'] < 0]['Experience'].value_counts()
# Convert Experience < 0 to positive number
Bank_Personal.loc[Bank_Personal['Experience']<0, 'Experience']=np.
↳abs(Bank_Personal['Experience'])
```

```
[14]: columns_to_drop = ['ID']
Bank_Personal = Bank_Personal.drop(columns=columns_to_drop)
Bank_Personal
```

```
[14]:
```

	Age	Experience	Income	ZIP Code	Family	CCAvg	Education	Mortgage	\
0	25	1	49	91107	4	1.6	1	0	
1	45	19	34	90089	3	1.5	1	0	
2	39	15	11	94720	1	1.0	1	0	
3	35	9	100	94112	1	2.7	2	0	
4	35	8	45	91330	4	1.0	2	0	
...	
4995	29	3	40	92697	1	1.9	3	0	
4996	30	4	15	92037	4	0.4	1	85	
4997	63	39	24	93023	2	0.3	3	0	

4998	65	40	49	90034	3	0.5	2	0
4999	28	4	83	92612	3	0.8	1	0

	Personal Loan	Securities Account	CD Account	Online	CreditCard
0	False	True	False	False	False
1	False	True	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	True
...
4995	False	False	False	True	False
4996	False	False	False	True	False
4997	False	False	False	False	False
4998	False	False	False	True	False
4999	False	False	False	True	True

[5000 rows x 13 columns]

There are different methods to identify outliers in the data. Some of these methods are:

Boxplot diagram: This method is used in the analysis of continuous data, and outliers can be identified by analyzing certain areas of the Boxplot diagram.

Cook's method: In this method, by calculating the average and standard deviation of the data, the permissible range for the data is determined, and then the data that are outside this range are known as outliers.

standard deviation method: In this method, data that are more than 3 standard deviations away from the average data are known as outliers. Boxplot diagram: This method is used in the analysis of continuous data, and by analyzing certain areas of the Boxplot diagram, outliers can be identified.

The method of the smallest and largest data: In this method, the value of the smallest and largest data is calculated, and then the data that are outside this range are known as outliers

```
[15]: #Treatment of outliers
from scipy import stats
for col in ['Mortgage', 'CCAvg']:
    # Calculate the mean and standard deviation of the data
    mean = Bank_Personal[col].mean()
    std = Bank_Personal[col].std()

    # Determining the threshold
    threshold = 3

    # Identifying outliers using z-score
    z_scores = (Bank_Personal[col] - mean) / std
    outliers = z_scores[abs(z_scores) > threshold].index
    print(f"number of outlier in {col} is :{Bank_Personal[stats.
↪zscore(Bank_Personal[col])>3][col].count()}")
```

```
## Remove outliers
Bank_Personal = Bank_Personal.drop(outliers)
Bank_Personal.reset_index(drop=True, inplace=True)
```

Bank_Personal

number of outlier in Mortgage is :105

number of outlier in CCAvg is :107

```
[15]:
```

	Age	Experience	Income	ZIP Code	Family	CCAvg	Education	Mortgage	\
0	25	1	49	91107	4	1.6	1	0	
1	45	19	34	90089	3	1.5	1	0	
2	39	15	11	94720	1	1.0	1	0	
3	35	9	100	94112	1	2.7	2	0	
4	35	8	45	91330	4	1.0	2	0	
...	
4783	29	3	40	92697	1	1.9	3	0	
4784	30	4	15	92037	4	0.4	1	85	
4785	63	39	24	93023	2	0.3	3	0	
4786	65	40	49	90034	3	0.5	2	0	
4787	28	4	83	92612	3	0.8	1	0	

	Personal Loan	Securities Account	CD Account	Online	CreditCard
0	False	True	False	False	False
1	False	True	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	True
...
4783	False	False	False	True	False
4784	False	False	False	True	False
4785	False	False	False	False	False
4786	False	False	False	True	False
4787	False	False	False	True	True

[4788 rows x 13 columns]

Attribute conversion

In this data set, CCAVG represents the average monthly credit card cost, but revenue represents the amount of annual revenue. To equalize the units of characteristics, we get the amount of monthly income

```
[16]: Bank_Personal['Income'] = Bank_Personal['Income']/12
```

```
[17]: cat_Features = ['CD Account', 'Education', 'Family', 'Securities_
↪Account', 'Online', 'CreditCard']
target_variable = 'Personal Loan'
```



```
[18]: from sklearn import preprocessing

# Loop over each column in the DataFrame where dtype is 'object'
for col in Bank_Personal.select_dtypes(include=['object','bool']).columns:

    # Print the column name and the unique values
    print(f"{col}: {Bank_Personal[col].unique()}")

# Loop over each column in the DataFrame where dtype is 'object'
for col in Bank_Personal.select_dtypes(include=['object','bool']).columns:
    # Convert all values to strings
    Bank_Personal[col] = Bank_Personal[col].astype(str)

    # Initialize a LabelEncoder object
    label_encoder = preprocessing.LabelEncoder()

    # Fit the encoder to the unique values in the column
    label_encoder.fit(Bank_Personal[col].unique())

    # Transform the column using the encoder
    Bank_Personal[col] = label_encoder.transform(Bank_Personal[col])

    # Print the column name and the unique encoded values
    print(f"{col}: {Bank_Personal[col].unique()}")
```

ZIP Code: [91107 90089 94720 94112 91330 92121 91711 93943 94710 90277 93106
94920

91741 95054 95010 94305 94015 90095 91320 95521 95064 90064 94539 94104
94117 94801 94035 92647 95814 94114 94115 92672 94122 95616 94065 95014
91380 95747 92373 92093 94005 90245 95819 90404 93407 94523 90024 91360
95670 95123 90045 91335 93907 92007 94606 94611 94901 92220 93305 95134
94612 92507 91730 94501 94303 94105 94550 92612 95617 92374 94080 94608
93555 93311 94704 92717 92037 95136 94542 94143 91775 92703 92354 92024
92831 92833 90057 92130 91301 92096 92646 92182 92131 90840 95035 93010
94928 95831 91770 94102 91423 93955 92834 93117 94551 94596 94025 94545
95053 90036 91125 95120 94706 95827 90503 90250 95817 93111 94132 95818
91942 90401 93524 95133 92173 94043 92521 92122 93118 92697 94577 91345
94123 92152 91355 94609 94306 96150 94110 94707 91604 90291 92807 95051
94085 92677 94304 92614 92626 94583 92103 92691 94107 92407 90504 94002
95039 94063 94923 95023 90058 92126 94118 90029 92806 94806 92110 94536
90623 93720 92069 92843 92120 90740 91207 95929 93437 90630 90034 90266
95630 92038 91304 92606 92192 90745 95060 94301 92692 92101 94610 94590
92028 92054 92029 93105 91941 92346 94402 94618 94904 9307 95482 91709
91311 94509 93023 92866 91745 90019 94111 94309 90073 92333 90505 94998
94086 94709 95825 90509 93108 94588 91706 94022 92109 92068 95841 92123
91342 90232 92634 91006 91768 90007 90028 92008 95112 92115 92177 90640
94607 92780 90009 92518 91007 93014 94024 90027 95207 90717 94534 94010

94234 90210 95020 92870 92124 90049 94521 95678 95045 92653 92821 90025
91910 94701 91129 95605 90071 96651 94960 91902 90033 95621 90037 90005
93940 91109 93009 93561 95126 94109 93107 94591 92251 92648 92709 91754
92009 96064 91103 91030 90066 95403 91016 95348 91950 95822 94538 91614
92154 92835 93657 93063 91040 92661 94061 95758 96091 90254 94066 94939
95138 95762 92064 94708 92106 92056 91302 90048 92325 91116 92868 90638
90747 93611 95833 91605 92675 90650 95820 90018 93711 95973 92886 95812
95503 91203 91105 95008 90016 90035 92129 90720 94949 90041 95003 95192
91101 94126 90230 93101 91365 91367 92660 92104 90405 91361 90011 90032
95354 94546 92673 95351 92399 90274 94087 90044 94131 94124 95032 90212
93109 94019 95828 90086 94555 93033 93022 91343 91911 94803 94553 95211
90304 92084 90601 92704 91763 92350 94705 93401 90502 94571 95070 92735
95037 95135 94028 96003 90065 95405 95370 93727 92867 95821 94566 95125
94526 94604 96008 93065 96001 95006 90639 92630 95307 91801 94302 91710
93950 90059 94558 93933 92161 94507 94575 95449 93403 93460 95005 93302
94040 91401 95816 92624 95131 94965 91784 91765 90280 95422 95518 95741
92694 90275 90272 94108 91791 92705 91773 93003 90755 96145 94703 96094
92116 95193 94116 90068 94970 90813 94404 94598 95842]

Personal Loan: [False True]

Securities Account: [True False]

CD Account: [False True]

Online: [False True]

CreditCard: [False True]

ZIP Code: [81 33 365 296 93 157 112 265 364 46 239 371 114 395 384 314 271
34


92 426 397 26 328 289 300 366 277 194 439 297 298 199 302 428 282 385
103 435 179 146 269 39 443 51 251 323 7 99 432 401 20 94 262 132
348 353 369 172 247 406 354 183 113 319 312 290 332 187 429 180 284 350
255 248 358 211 138 408 329 308 122 207 178 135 217 218 23 163 88 147
193 170 164 72 389 229 373 450 120 288 105 267 219 244 333 345 275 330
394 16 84 400 360 448 54 40 441 243 307 442 130 50 254 405 168 279
185 158 245 206 340 97 303 165 98 351 315 463 294 361 106 48 215 393
285 202 313 188 190 341 149 203 291 182 55 268 391 281 372 387 24 161
301 11 214 368 153 326 59 260 144 221 156 67 87 454 252 60 14 42
431 139 90 186 171 68 396 310 204 148 352 343 136 140 137 238 129 176
317 355 370 236 423 110 91 321 232 222 115 6 295 316 31 175 56 379
286 363 447 57 241 342 109 273 152 143 452 159 95 38 192 73 119 1
10 133 399 154 169 63 349 213 2 184 74 230 274 9 412 65 325 270
309 35 386 225 160 22 322 433 392 196 216 8 127 356 85 427 30 464
376 126 13 430 17 0 264 82 228 256 403 293 240 344 173 195 210 116
134 459 79 76 28 419 75 415 131 446 327 108 166 220 258 234 77 198
280 436 460 41 283 374 409 437 142 362 151 141 89 21 174 83 224 61
69 257 451 107 201 64 444 5 259 455 226 438 424 86 80 383 4 15
162 66 375 18 380 410 78 305 37 237 101 102 197 150 52 100 3 12
417 331 200 416 181 44 287 19 306 304 388 36 242 272 449 32 335 233
231 96 128 367 334 413 49 145 58 208 117 177 359 249 53 338 398 212
390 407 276 457 27 420 418 261 223 445 337 402 324 347 458 235 456 382
62 191 414 125 311 111 266 25 336 263 167 320 339 422 250 253 381 246

```
278 104 440 189 404 377 123 118 47 421 425 434 205 45 43 292 124 209
121 227 70 462 357 461 155 411 299 29 378 71 318 346 453]
Personal Loan: [0 1]
Securities Account: [1 0]
CD Account: [0 1]
Online: [0 1]
CreditCard: [0 1]
```

#Correlation

If the correlation coefficient between two variables is greater than 0.7, it is considered as a strong correlation. If the correlation coefficient between two variables is between 0.3 and 0.7, it is considered as a moderate correlation. If the correlation coefficient between two variables is less than 0.3, it is considered as a weak correlation.

```
[19]: fig = plt.figure(figsize=(12,12))
sns.heatmap(Bank_Personal.corr(),cmap='Blues',annot =True)
plt.savefig("correlation")
plt.show()
```



Personal Loan is highly correlated with Income, CD_Account, CCAvg.

Experience is highly correlated with Age. ($r = 1$) CCAvg is correlated with Income to a good extent. ($r = 0.6$)

Age and Experience features have very high correlation, 0.99. It is also intuitively understandable that experience increases as age increases. Correlated features degrade the learning performance and causes instability on the models

#Data Analysis

sns.pairplot is a powerful tool for getting a visual overview of the relationships and distributions within a dataset. It is especially useful for exploratory data analysis (EDA) to identify patterns,

trends, and potential areas of interest for further investigation. And also, it is a powerful tool for getting a visual overview of the relationships and distributions within a dataset.

Univariate Distributions:

Along the diagonal of the pairplot, you'll see histograms or kernel density plots for each variable, showing the univariate distribution of each variable.

Bivariate Relationships:

In the lower and upper triangles of the pairplot grid, scatterplots are displayed for pairs of variables, showing the bivariate relationships between them. Each point in the scatterplot represents an observation in the dataset.

Correlation Information:

By visually inspecting the scatterplots, you can get an idea of the correlation or relationship between different pairs of variables. The orientation and shape of the scatterplots can indicate whether there is a positive or negative correlation.

Identifying Outliers:

Outliers or unusual observations in the data can often be visually identified in the scatterplots. Points that deviate significantly from the overall pattern may stand out.

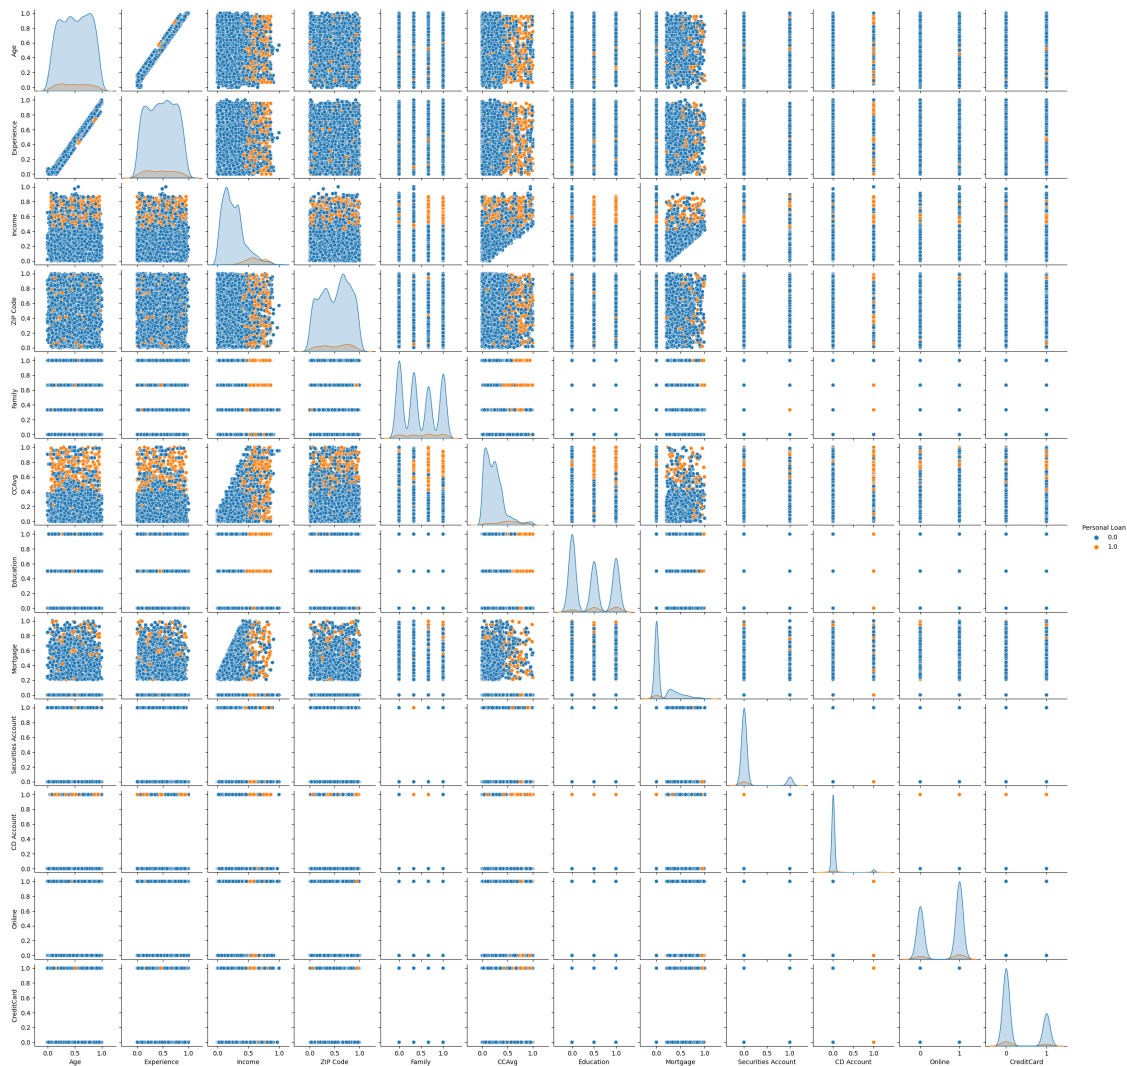
Pairwise Comparisons:

The pairplot allows you to quickly compare all pairs of variables in the dataset, making it useful for gaining an overview of relationships and potential patterns.

Faceting by Categorical Variables:

If the dataset contains categorical variables, you can use the hue parameter to create separate subplots for different categories. This allows you to see how the relationships differ across categories.

```
[ ]: %matplotlib inline
import seaborn as sns
sns.pairplot(Bank_Personal, hue='Personal Loan', height=2);
plt.savefig('pairplot.png')
```





By examining both the histogram and boxplot for each column, we can gain insights into the distribution, skewness, and presence of outliers in the data.


This is useful for visualizing the univariate distribution of each column, helping you identify patterns and potential issues in the data

```
[20]: for col in Bank_Personal.columns:
    fig, axs = plt.subplots(nrows=2, figsize=(7,7))
    print("column name :",col)
    axs[0].hist(Bank_Personal[col], bins=20)
    axs[1].boxplot(Bank_Personal[col])
    plt.tight_layout()
    plt.show()
```


column name : Age




column name : Experience




column name : Income




column name : ZIP Code




column name : Family




column name : CCAvg




column name : Education




column name : Mortgage




column name : Personal Loan




column name : Securities Account




column name : CD Account



column name : Online



column name : CreditCard



The code generates a grid of count plots for each categorical feature in our DataFrame against the target variable(personal loan) Information Provided:

Comparison of Categorical Features:

Each subplot shows the count distribution of categories within a specific categorical feature. Differences in counts can be observed across categories for each feature.

Impact on Target Variable:

The hue (color) represents the target variable, allowing us to visually assess how the distribution of the target variable varies across categories of each feature.

Insights into Relationships:

Provides insights into how each categorical feature might be related to the target variable. Useful

for identifying patterns and making decisions about feature importance.

```
[35]: fig=plt.figure(figsize=(15,10))
      for i,col in enumerate(cat_Features):
          ax=fig.add_subplot(2,4,i+1)
          sns.barplot(x=col,y='Personal Loan',data=Bank_Personal,ci=None)
          plt.savefig("subplotxx")
```

<ipython-input-35-5ff7ddfcd77>:4: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(x=col,y='Personal Loan',data=Bank_Personal,ci=None)
```

<ipython-input-35-5ff7ddfcd77>:4: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(x=col,y='Personal Loan',data=Bank_Personal,ci=None)
```

<ipython-input-35-5ff7ddfcd77>:4: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(x=col,y='Personal Loan',data=Bank_Personal,ci=None)
```

<ipython-input-35-5ff7ddfcd77>:4: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(x=col,y='Personal Loan',data=Bank_Personal,ci=None)
```

<ipython-input-35-5ff7ddfcd77>:4: FutureWarning:


The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(x=col,y='Personal Loan',data=Bank_Personal,ci=None)
```

<ipython-input-35-5ff7ddfcd77>:4: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.


```
sns.barplot(x=col,y='Personal Loan',data=Bank_Personal,ci=None)
```



Customers with family size equal to 3 have more chances of having Personal Loan. Customers with Undergraduate degree have less chances of having Personal Loan as compared to other customers having Graduate or Advanced/Professional degree. Customers with CD Account and Securities Account have more chances of having Personal Loan. Customers with Online & Credit Card is more likely to have Personal Loan than others don't have a one.

```
[24]: sns.kdeplot(
    data=Bank_Personal, x='Income', hue="Personal Loan",
    fill=True, common_norm=False, palette="husl",
    alpha=.5, linewidth=0,
)
```

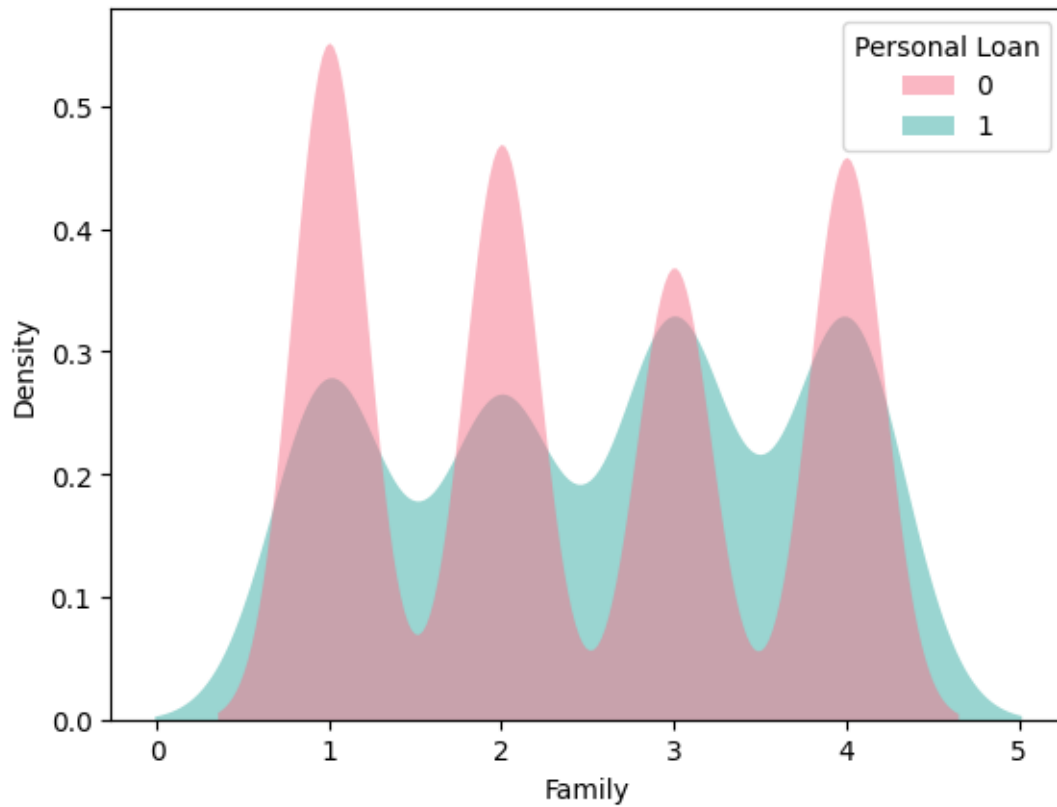
```
[24]: <Axes: xlabel='Income', ylabel='Density'>
```



the chart shows that Personal Loan is more likely to happen when the Income increases.

```
[25]: sns.kdeplot(  
    data=Bank_Personal, x='Family', hue="Personal Loan",  
    fill=True, common_norm=False, palette="husl",  
    alpha=.5, linewidth=0,  
)
```

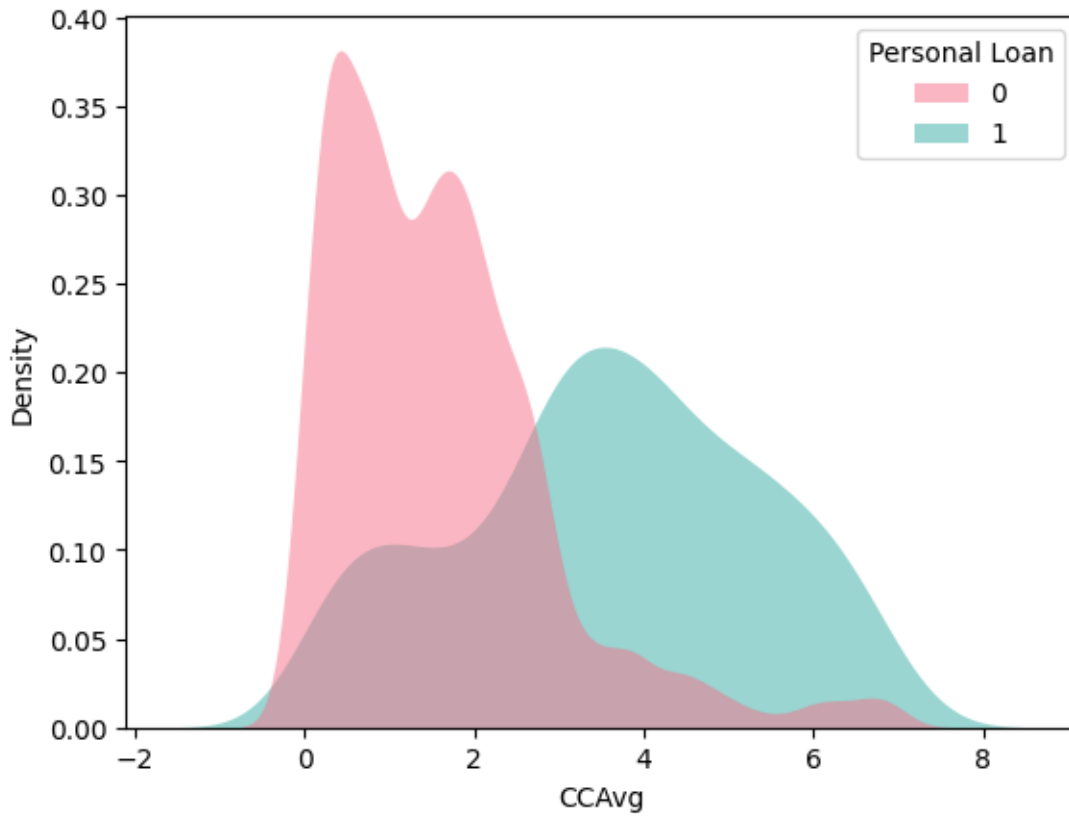
```
[25]: <Axes: xlabel='Family', ylabel='Density'>
```



the chart shows that Personal Loan is more likely to be achieved when the Family members are bigger.

```
[29]: sns.kdeplot(  
    data=Bank_Personal, x='CCAvg', hue="Personal Loan",  
    fill=True, common_norm=False, palette="husl",  
    alpha=.5, linewidth=0,  
)
```


```
[29]: <Axes: xlabel='CCAvg', ylabel='Density'>
```



the chart shows that Personal Loan is more likely to happen when the Credit Card Average of spending is higher.

```
[30]: sns.kdeplot(  
    data=Bank_Personal, x='Education', hue="Personal Loan",  
    fill=True, common_norm=False, palette="husl",  
    alpha=.5, linewidth=0,  
)
```


```
[30]: <Axes: xlabel='Education', ylabel='Density'>
```

the chart shows that Personal Loan is more likely to happen when the Education level is higher.

```
[31]: sns.kdeplot(
      data=Bank_Personal, x='Mortgage', hue="Personal Loan",
      fill=True, common_norm=False, palette="husl",
      alpha=.5, linewidth=0,
    )
```

```
[31]: <Axes: xlabel='Mortgage', ylabel='Density'>
```




the chart shows that Personal Loan is more likely to happen when the Mortgage equal 0

```
[36]: plt.figure(figsize=(35,20),dpi=90)
      # Plot frequency percentages barplot
      Bank_Personal['Personal Loan'].value_counts(normalize=True).mul(100).
      ↪plot(kind='barh', width=0.8, figsize=(8,5))

      # Add frequency percentages to the plot
      labels = Bank_Personal['Personal Loan'].value_counts(normalize=True).mul(100).
      ↪round(1)
      for i in labels.index:
          plt.text(labels[i], i, str(labels[i])+ '%', fontsize=15, weight='bold')

      plt.xlim([0, 110])
      plt.xlabel('Frequency Percentage', fontsize=15)
      plt.ylabel('Personal Loan', fontsize=15)
      plt.title('Frequency Percentage of Target Classes', fontsize=15)
      plt.show()
```



This code generates a scatter plot using Seaborn, visualizing the relationship between two numerical variables, 'Income' and 'CCAvg', with points color-coded based on the 'Personal Loan' variable. This is helpful for understanding the correlation or patterns between two numerical variables. Allows exploration of how the 'Personal Loan' variable relates to the 'Income' and 'CCAvg' dimensions.

```
[38]: sns.scatterplot(x='Income', y='CCAvg', hue='Personal Loan', data=Bank_Personal,
    ↪alpha=0.5, palette='husl')
```

```
[38]: <Axes: xlabel='Income', ylabel='CCAvg'>
```



#Applying scaling

Min-max scaling is beneficial when numerical features have different scales, ensuring that each feature contributes equally to the model. The code is useful for preprocessing data, especially when preparing it for machine learning algorithms that are sensitive to feature scales.

Note:

Min-max scaling transforms the data such that the minimum value becomes 0, the maximum value becomes 1, and values in between are linearly scaled. The transformation is column-wise, and each feature is scaled independently

```
[39]: from sklearn.preprocessing import MinMaxScaler
numerical_columns = [col for col in Bank_Personal.
    ↪select_dtypes(include=['float64', 'int64'])]
# Initialize MinMaxScaler
scaler = MinMaxScaler()
# Fit and transform the numerical columns for Test_space (excluding
    ↪'PassengerId')
Bank_Personal[numerical_columns] = scaler.
    ↪fit_transform(Bank_Personal[numerical_columns])
Bank_Personal
```

```
[39]:
```

	Age	Experience	Income	ZIP Code	Family	CCAvg	Education \
0	0.045455	0.023256	0.189815	0.174569	1.000000	0.228571	0.0
1	0.500000	0.441860	0.120370	0.071121	0.666667	0.214286	0.0
2	0.363636	0.348837	0.013889	0.786638	0.000000	0.142857	0.0
3	0.272727	0.209302	0.425926	0.637931	0.000000	0.385714	0.5
4	0.272727	0.186047	0.171296	0.200431	1.000000	0.142857	0.5
...
4783	0.136364	0.069767	0.148148	0.443966	0.000000	0.271429	1.0
4784	0.159091	0.093023	0.032407	0.297414	1.000000	0.057143	0.0
4785	0.909091	0.906977	0.074074	0.500000	0.333333	0.042857	1.0
4786	0.954545	0.930233	0.189815	0.030172	0.666667	0.071429	0.5
4787	0.113636	0.093023	0.347222	0.403017	0.666667	0.114286	0.0

	Mortgage	Personal Loan	Securities Account	CD Account	Online \
0	0.000000	0.0	1.0	0.0	0.0
1	0.000000	0.0	1.0	0.0	0.0
2	0.000000	0.0	0.0	0.0	0.0
3	0.000000	0.0	0.0	0.0	0.0
4	0.000000	0.0	0.0	0.0	0.0
...
4783	0.000000	0.0	0.0	0.0	1.0
4784	0.235457	0.0	0.0	0.0	1.0
4785	0.000000	0.0	0.0	0.0	0.0
4786	0.000000	0.0	0.0	0.0	1.0
4787	0.000000	0.0	0.0	0.0	1.0

	CreditCard
0	0.0
1	0.0
2	0.0
3	0.0
4	1.0
...	...
4783	0.0
4784	0.0
4785	0.0
4786	0.0
4787	1.0

```
[4788 rows x 13 columns]
```

```
#Train-Test Splitting
```


```
[40]: from sklearn.model_selection import train_test_split, StratifiedKFold,
      ↪ cross_val_score, cross_validate, GridSearchCV
      from sklearn.metrics import accuracy_score, precision_score, recall_score,
      ↪ f1_score, classification_report, confusion_matrix, roc_auc_score, make_scorer
```

```
[41]: x_o=Bank_Personal.drop('Personal Loan',axis=1)
      y_o=Bank_Personal['Personal Loan']
```

```
[42]: from sklearn.model_selection import train_test_split
      x_train_o,x_test_o,y_train_o,y_test_o=train_test_split(x_o,y_o,test_size=0.
      ↪2,random_state=42,stratify=y_o)
```

```
[43]: mutual_information = mutual_info_classif(x_train_o, y_train_o, n_neighbors=5,
      ↪copy = True)

      plt.subplots(1, figsize=(26, 1))
      sns.heatmap(mutual_information[:, np.newaxis].T, cmap='Blues', cbar=False,
      ↪linewidths=1, annot=True, annot_kws={"size": 20})
      plt.yticks([], [])
      plt.gca().set_xticklabels(x_train_o.columns, rotation=45, ha='right',
      ↪fontsize=16)
      plt.suptitle("Variable Importance (mutual_info_classif)", fontsize=22, y=1.2)
      plt.gcf().subplots_adjust(wspace=0.2)
```



Observation: most Important features on dataset (Income, CCAvg, CD Account)

```
[44]: y_o.value_counts()
```

```
[44]: 0.0    4381
      1.0     407
      Name: Personal Loan, dtype: int64
```

```
[45]: y_train_o.value_counts()
```

```
[45]: 0.0    3504
      1.0     326
      Name: Personal Loan, dtype: int64
```

```
[46]: y_test_o.value_counts()
```

```
[46]: 0.0     877
      1.0      81
      Name: Personal Loan, dtype: int64
```

#Models

first we want to check what kind of model is better for us for this classification. and we check the accuracy and confusion matrix as an parameter to find the best model for classification. as we find out the tree classifier and random forest has the most accuracy so we choose them. but we just check naive bayes. based on our curocity to work on that and check the performnce.

```
[47]: from sklearn.model_selection import KFold
      from sklearn.model_selection import cross_val_score
      from sklearn.linear_model import LinearRegression
      from sklearn.linear_model import LogisticRegression
      from sklearn.model_selection import GridSearchCV
      from sklearn.metrics import classification_report
      from sklearn import svm, tree
      from sklearn.ensemble import AdaBoostClassifier, RandomForestClassifier

      random_state = 42
      train_s = 20
      cv = KFold(n_splits=10, random_state=1, shuffle=True)

      classifiers = []
      model1 = svm.SVC()
      classifiers.append(model1)
      model2 = tree.DecisionTreeClassifier()
      classifiers.append(model2)
      model3 = RandomForestClassifier()
      classifiers.append(model3)
      model4 = LogisticRegression(random_state=1)
      classifiers.append(model4)
      model5 = GaussianNB()
      classifiers.append(model5)
```

```
[48]: for clf in classifiers:
      clf.fit(x_train_o, y_train_o)
      y_pred= clf.predict(x_test_o)
      acc = accuracy_score(y_test_o, y_pred)
      print("Accuracy of %s is %s"%(clf, acc))
      cm = confusion_matrix(y_test_o, y_pred)
      print("Confusion Matrix of %s is %s"%(clf, cm))
```

```
Accuracy of SVC() is 0.9718162839248434
Confusion Matrix of SVC() is [[873   4]
 [ 23  58]]
Accuracy of DecisionTreeClassifier() is 0.9812108559498957
Confusion Matrix of DecisionTreeClassifier() is [[869   8]
 [ 10  71]]
Accuracy of RandomForestClassifier() is 0.9895615866388309
Confusion Matrix of RandomForestClassifier() is [[876   1]
 [   9  72]]
Accuracy of LogisticRegression(random_state=1) is 0.954070981210856
Confusion Matrix of LogisticRegression(random_state=1) is [[865  12]
```

```
[ 32  49]]
Accuracy of GaussianNB() is 0.894572025052192
Confusion Matrix of GaussianNB() is [[812  65]
[ 36  45]]
```

#Decision Tree Model

```
[49]: estimator = tree.DecisionTreeClassifier(criterion="entropy", random_state = 42)
estimator.fit(x_train_o, y_train_o);
```

```
[50]: y_predicted_train = estimator.predict(x_train_o)
accuracy_train = accuracy_score(y_predicted_train, y_train_o)*100
print("The accuracy on training set is {0:.1f}%".format(accuracy_train))
```

The accuracy on training set is 100.0%

```
[51]: y_predicted_test = estimator.predict(x_test_o)
accuracy_ho = accuracy_score(y_test_o, y_predicted_test) * 100
fitted_max_depth = estimator.tree_.max_depth
fitted_x = train_s
initial_impurity = estimator.tree_.impurity[0] # the impurity variable of tree_
↳contains the impurities of all the nodes
print("The accuracy on test set is {0:.1f}%".format(accuracy_ho))
print("The maximum depth of the tree fitted on X_train is {}".
      ↳format(fitted_max_depth))

parameter_values = range(1,fitted_max_depth+1)
# parameter_values = np.linspace(0, initial_impurity, 31)
```

The accuracy on test set is 98.5%


The maximum depth of the tree fitted on X_train is 14

```
[52]: avg_scores = []
for par in parameter_values:
    estimator = tree.DecisionTreeClassifier(criterion="entropy"
                                             , max_depth = par
                                             , random_state = random_state
                                             )
    scores = cross_val_score(estimator, x_train_o, y_train_o
                             , scoring='accuracy', cv = 5)
    # cross_val_score produces an array with one score for each fold
    avg_scores.append(np.mean(scores))
print(avg_scores)
```

```
[0.9148825065274151, 0.9514360313315928, 0.9783289817232376, 0.985378590078329,
0.9843342036553524, 0.9830287206266318, 0.9796344647519583, 0.9825065274151437,
0.9814621409921671, 0.9817232375979111, 0.9819843342036554, 0.9822454308093993,
0.9832898172323759, 0.9832898172323759]
```



```
[53]: plt.figure(figsize=(32,20))
plt.plot(parameter_values, avg_scores, '-o', linewidth=5, markersize=24)
plt.xlabel('max_depth')
plt.ylabel('accuracy')
plt.title("Score with Cross Validation varying max_depth of tree", fontsize = 24)
plt.show();
```



```
[54]: top_par_cv = parameter_values[np.argmax(avg_scores)]
estimator = tree.DecisionTreeClassifier(criterion="entropy", max_depth = top_par_cv)
estimator.fit(x_train_o,y_train_o);
y_predicted = estimator.predict(x_test_o)
accuracy_cv = accuracy_score(y_test_o, y_predicted) * 100
print("The accuracy on test set tuned with cross_validation is {:.1f}% with depth of the tree {}".format(accuracy_cv , top_par_cv))
```

The accuracy on test set tuned with cross_validation is 98.5% with depth of the tree 4

```
[ ]: print(classification_report(y_test_o, y_predicted))
```

```
precision    recall  f1-score   support
```

0.0	0.99	0.99	0.99	877
1.0	0.94	0.89	0.91	81
accuracy			0.99	958
macro avg	0.96	0.94	0.95	958
weighted avg	0.99	0.99	0.99	958

Class 0 (0.0):

Precision: 1.00

Interpretation: 100% of the instances predicted as Class 0 were actually Class 0.

Recall (Sensitivity or True Positive Rate): 1.00

Interpretation: The model correctly identified 100% of the instances of Class 0.

F1-Score: 1.00

Interpretation: The balance between precision and recall for Class 0 is excellent.

Class 1 (1.0):

Precision: 0.97

Interpretation: 97% of the instances predicted as Class 1 were actually Class 1.

Recall: 0.96


Interpretation: The model identified 96% of the instances of Class 1.

F1-Score: 0.96

Interpretation: The balance between precision and recall for Class 1 is very good.

The high precision and recall for both classes suggest that the model is performing very well on the given dataset. The F1-score provides a balanced measure of performance, considering both precision and recall.

```
[ ]: cm = confusion_matrix(y_test_o, y_predicted)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=['Not has',
↳Personal Loan', 'Personal Loan'], yticklabels=['Not has Personal Loan',
↳Personal Loan'])
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```



#Random Forest Model

```
[ ]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV
# Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 1, stop = 250, num = 1)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Method of selecting samples for training each tree
bootstrap = [True, False]
# Create the random grid
```

```

random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap}

print(random_grid)

```

```

{'n_estimators': [1], 'max_features': ['auto', 'sqrt'], 'max_depth': [10, 20,
30, 40, 50, 60, 70, 80, 90, 100, 110, None], 'min_samples_split': [2, 5, 10],
'min_samples_leaf': [1, 2, 4], 'bootstrap': [True, False]}

```

```

[ ]: rf = RandomForestClassifier()
rf_random = RandomizedSearchCV(estimator = rf, param_distributions =_
    random_grid, n_iter = 100, cv = 3, verbose=2, random_state=42, n_jobs = -1)
rf_random.fit(x_train_o,y_train_o)

```

Fitting 3 folds for each of 100 candidates, totalling 300 fits

```

[ ]: RandomizedSearchCV(cv=3, estimator=RandomForestClassifier(), n_iter=100,
    n_jobs=-1,
    param_distributions={'bootstrap': [True, False],
                        'max_depth': [10, 20, 30, 40, 50, 60,
                                     70, 80, 90, 100, 110,
                                     None],
                        'max_features': ['auto', 'sqrt'],
                        'min_samples_leaf': [1, 2, 4],
                        'min_samples_split': [2, 5, 10],
                        'n_estimators': [1]}},
    random_state=42, verbose=2)

```

```

[ ]: randomforestmodel=RandomForestClassifier(n_estimators= 1,
    min_samples_split = 10,
    min_samples_leaf = 2,
    max_features = 'sqrt',
    max_depth= 90,
    bootstrap= False)

```

```

[ ]: randomforestmodel.fit(x_train_o,y_train_o)

```

```

[ ]: RandomForestClassifier(bootstrap=False, max_depth=90, min_samples_leaf=2,
    min_samples_split=10, n_estimators=1)

```

```

[ ]: y_pred_rf=randomforestmodel.predict(x_test_o)

```

```

[ ]: rf_accuracy = accuracy_score(y_test_o, y_pred_rf)
rf_conf_matrix = confusion_matrix(y_test_o, y_pred_rf)

```

```
rf_classification_rep = classification_report(y_test_o, y_pred_rf)
```

```
[ ]: print(f'rf_Accuracy: {rf_accuracy:.2f}')
```

```
rf_Accuracy: 0.98
```

```
[ ]: print(rf_classification_rep)
```

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	877
1.0	0.95	0.85	0.90	81
accuracy			0.98	958
macro avg	0.97	0.92	0.94	958
weighted avg	0.98	0.98	0.98	958

The Random Forest classifier performs well for both classes, with high precision, recall, and F1-score for Class 0. For Class 1, the model still maintains a good balance between precision and recall, but the metrics are slightly lower compared to Class 0.

Precision: 0.98

Interpretation: 98% of the instances predicted as Class 0 were actually Class 0.

Recall (Sensitivity or True Positive Rate): 0.98

Interpretation: The model correctly identified 98% of the instances of Class 0.


F1-Score: 0.98

Interpretation: The balance between precision and recall for Class 0 is excellent.

Interpretation:

High precision indicates that when the model predicts a certain class, it is likely to be correct. High recall indicates that the model can capture a large proportion of actual instances of a class. F1-score is a harmonic mean of precision and recall, providing a balance between the two metrics.

```
[ ]: cm = rf_conf_matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=['Not has',
    ↳Personal Loan', 'Personal Loan'], yticklabels=['Not has Personal Loan',
    ↳Personal Loan'])
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```



#Naive Base Model

```
[ ]: GaussianClassifier = GaussianNB()
GaussianClassifier.fit(x_train_o,y_train_o)
y_pred=GaussianClassifier.predict(x_test_o)
print("=====> Result_")
    <=====")
print("Accuracy          = " ,metrics.
    <accuracy_score(y_test_o,y_pred))
print("F1 Score          = " ,metrics.f1_score(y_test_o,y_pred))

=====> Result <=====
Accuracy          = 0.894572025052192
F1 Score          = 0.4712041884816754

[ ]: cm = confusion_matrix(y_test_o,GaussianClassifier.predict(x_test_o))

plt.figure(figsize=(8, 6))
```

```

sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=['Not has_
↳Personal Loan', 'Personal Loan'], yticklabels=['Not has Personal Loan',_
↳'Personal Loan'])
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()

```



```
[ ]: print(classification_report(y_o,GussianClassifier.predict(x_o)))
```

	precision	recall	f1-score	support
0.0	0.96	0.93	0.95	4381
1.0	0.45	0.60	0.51	407
accuracy			0.90	4788
macro avg	0.70	0.77	0.73	4788
weighted avg	0.92	0.90	0.91	4788

Class 0 (0.0):

Precision: 0.95

Interpretation: 95% of the instances predicted as Class 0 were actually Class 0.

Recall (Sensitivity or True Positive Rate): 0.91

Interpretation: The model correctly identified 91% of the instances of Class 0.

F1-Score: 0.93

Interpretation: The balance between precision and recall for Class 0 is good.

Class 1 (1.0):

Precision: 0.41

Interpretation: 41% of the instances predicted as Class 1 were actually Class 1.

Recall: 0.57

Interpretation: The model identified 57% of the instances of Class 1.

F1-Score: 0.48

Interpretation: The balance between precision and recall for Class 1 is moderate.

Interpretation:

High precision for Class 0 indicates that when the model predicts Class 0, it is likely to be correct. The lower precision for Class 1 suggests that the model may have some false positives for this class. The recall for Class 1 suggests that the model captures more than half of the actual instances of Class 1. The F1-score provides a balanced measure of performance, considering both precision and recall.